

A 3D visualization of a distributed data structure. It features a grid of light gray rectangular blocks arranged in a roughly square pattern. Red lines are drawn across the top surfaces of these blocks, forming a complex, interconnected network that represents data flow or a shared log structure. The perspective is from an elevated angle, looking down at the grid.

Tango: Distributed Data Structures over a Shared Log

Group 10

- André Matos 92420
- Diogo Ribeiro 102484
- Luís Calado 103883

Index

- Introduction
 - Shared Log
- Tango objects
 - Transactions
- Layered partitions
- Transactions over streams
 - Evaluation
 - Conclusion

Shared Log

Tango builds on the CORFU shared log abstraction

Log uses chain replication

CORFU operations:

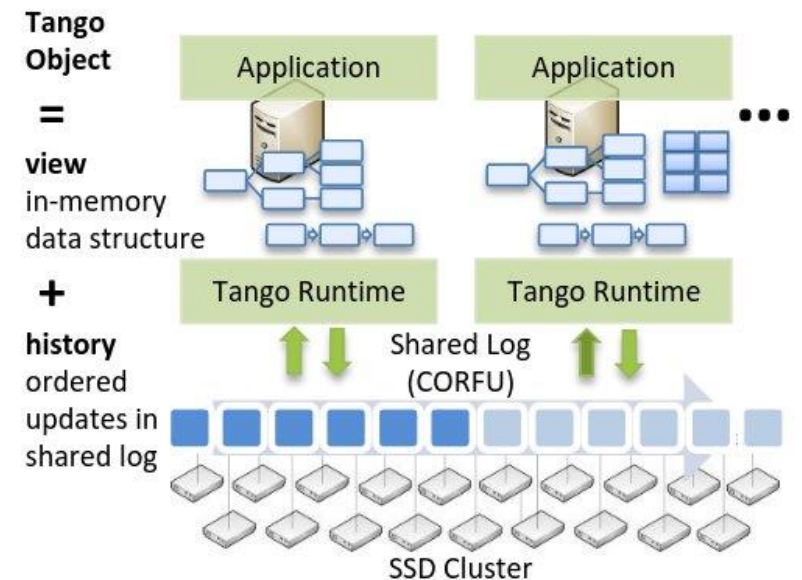
- **append(b)** -> append entry b to the shared log, obtaining an offset/position l in return
- **check()** -> check the current tail of the log
- **read(l)** -> read the entry at log position l
- **trim(l)** -> trim a particular offset in the log for garbage collection

No communication between clients.

CORFU Sequencer is used.

The Tango Object = View + History

- **Structure:** Historical log of updates + multiple views (tree, map).
- **Client Interaction:** Append updates to log, sync views with apply upcall.
- **Custom Views:** Different views (tree, set) from same log.
- **Consistency & Durability:** Linearizability, recovery via log history, checkpoints.



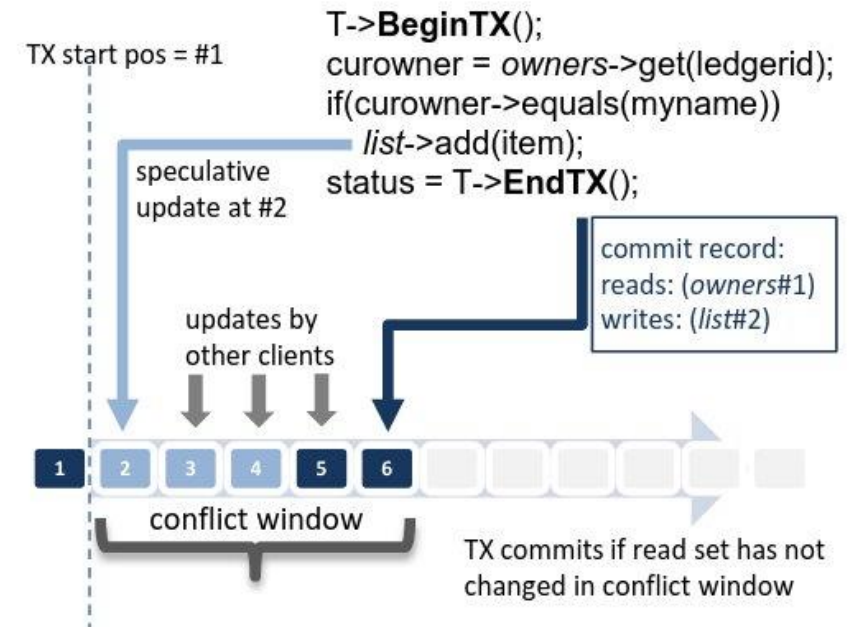
TangoRegister

- **update helper**: this accepts an opaque buffer from the object and appends it to the shared log.
- **query helper**: this reads new entries from the shared log and provides them to the object via an apply upcall.

```
class TangoRegister {  
    int oid;  
    TangoRuntime *T;  
    int state;  
    void apply(void *X) {  
        state = *(int *)X;  
    }  
    void writeRegister(int newstate){  
        T->update_helper(&newstate ,  
                        sizeof(int), oid);  
    }  
    int readRegister() {  
        T->query_helper(oid);  
        return state;  
    }  
}
```

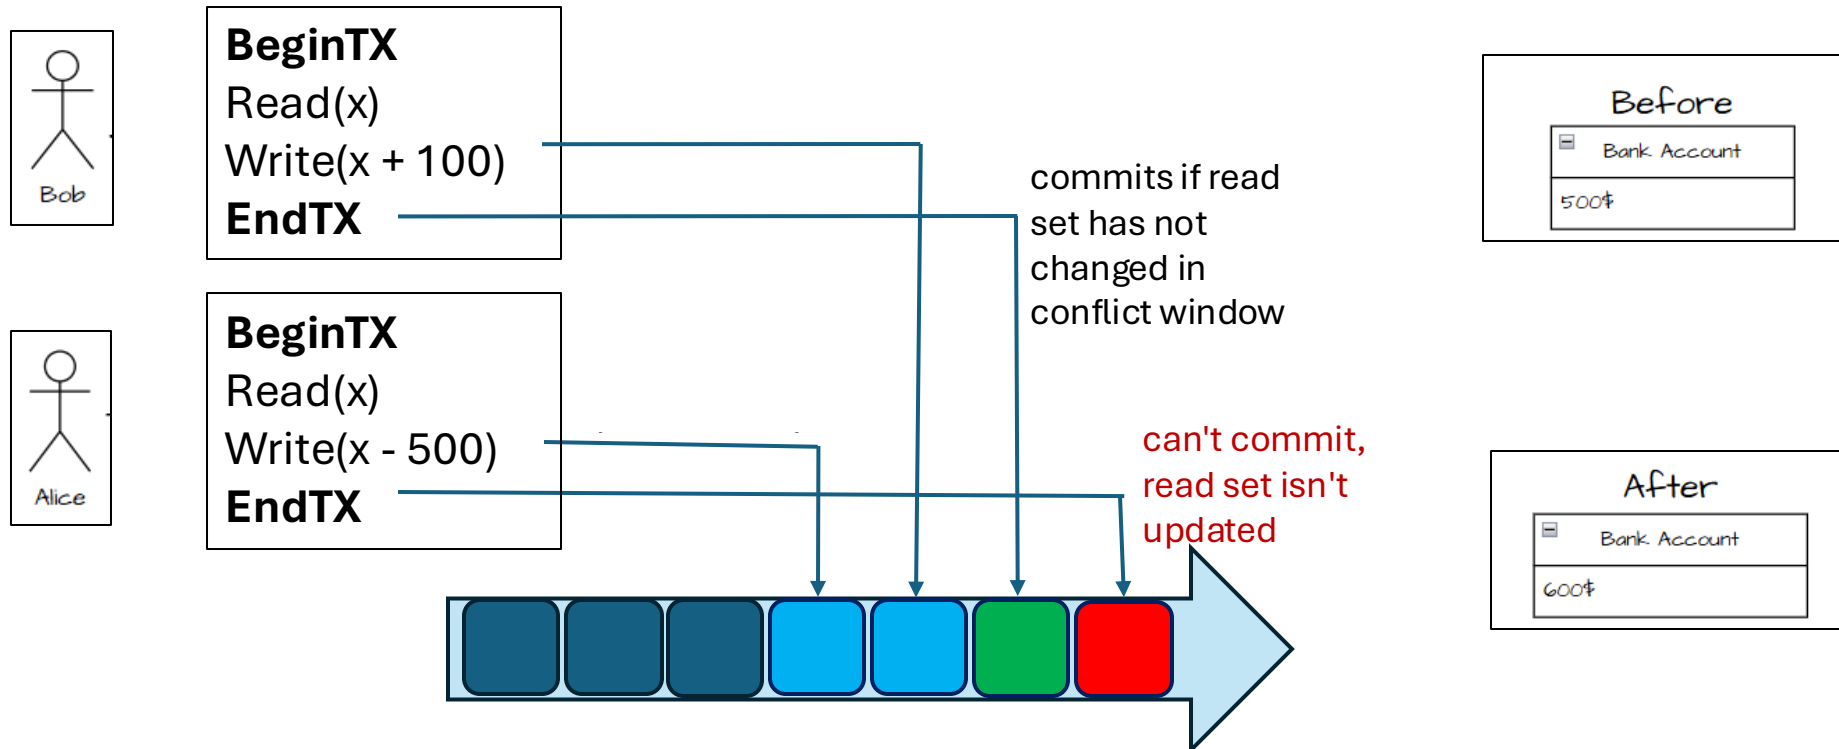
Transactions

- **Concurrency:** Speculative commit, decision after log check.
- **Commit Record:** Tracks readset (objects/versions), checks for changes.
- **Atomicity/Isolation:** Single log ensures multi-object transactions.
- **Transaction Flow:** BeginTX sets context, EndTX commits or aborts.



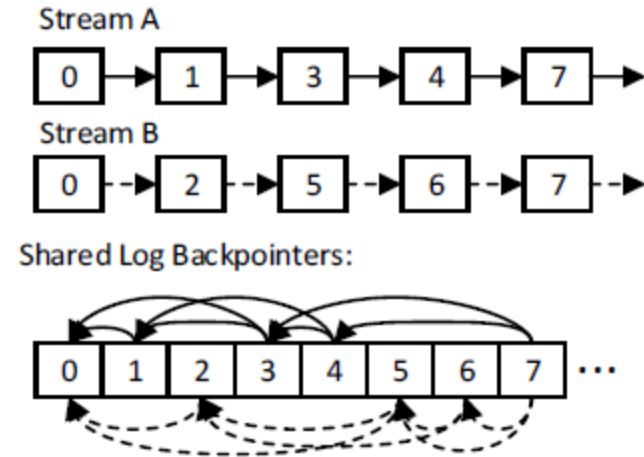


Transactions

- **Read/Write:** Read-only skips commit; write-only skips log traversal.



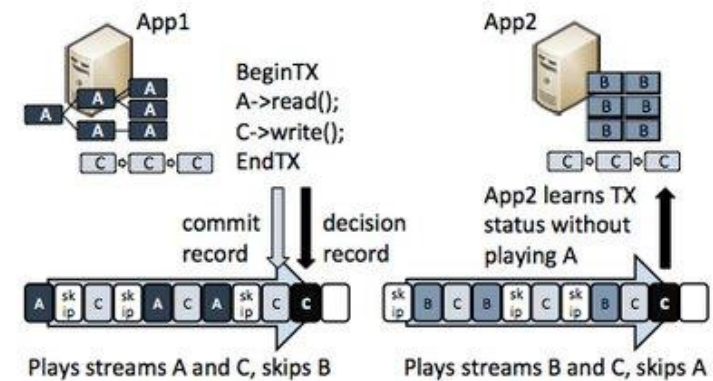

Layered Partitions

- **Partitioning:** Objects mapped to streams; clients read relevant streams.
- **Streams:** Linked lists built through backpointers.
- **Append:** Stream IDs sent to sequencer, respond with backpointers.
- **Failure Handling:** Reverse log traversal, junk values fill gaps.




Transactions Over Streams

- **Multiappend:** Transactions affect multiple streams, single commit in shared log.
- **Limitation:** Can't involve reads from objects without local copies.
- **Commit Decision:** Client adds commit decision to log; other clients can add if needed.
- **Failure Handling:** If commit fails, any client can reconstruct objects from the log after a timeout.




Evaluation





Two views

Evaluation



Evaluation



Conclusion

Main Contributions:

- Different structures, same log
- Clients with different views
- Clients choose views

